

Today's Reading, Pavlo 2009

MR VS. DATABASES


Hadoop vs. RDBMS

- Comparison of 3 systems
 - Hadoop
 - Vertica (a column-oriented database)
 - DBMS-X (a row-oriented database)
 - rhymes with “schmoracle”
- Qualitative
 - Programming model, ease of setup, features, etc.
- Quantitative
 - Data loading, different types of queries


Grep Task

- Find 3-byte pattern in 100-byte record
 - *1 match per 10,000 records*
- Data set:
 - *10-byte unique key, 90-byte value*
 - *1TB spread across 25, 50, or 100 nodes*
 - *10 billion records*
- Original MR Paper (Dean et al. 2004)

Grep Task Loading Results




Grep Task Execution Results



Selection Task

```
SELECT pageURL, pageRank  
FROM Rankings WHERE pageRank > X
```



1 GB /
node

Analytical Tasks


- Simple web processing schema
- Data set:
 - 600k HTML Documents (6GB/node)
 - 155 million UserVisit records (20GB/node)
 - 18 million Rankings records (1GB/node)

Aggregate Task

- Simple query to find adRevenue by IP prefix

```
SELECT SUBSTR(sourceIP, 1, 7),  
       SUM(adRevenue)  
FROM   userVistits  
GROUP BY SUBSTR(sourceIP, 1, 7)
```


Aggregate Task Results



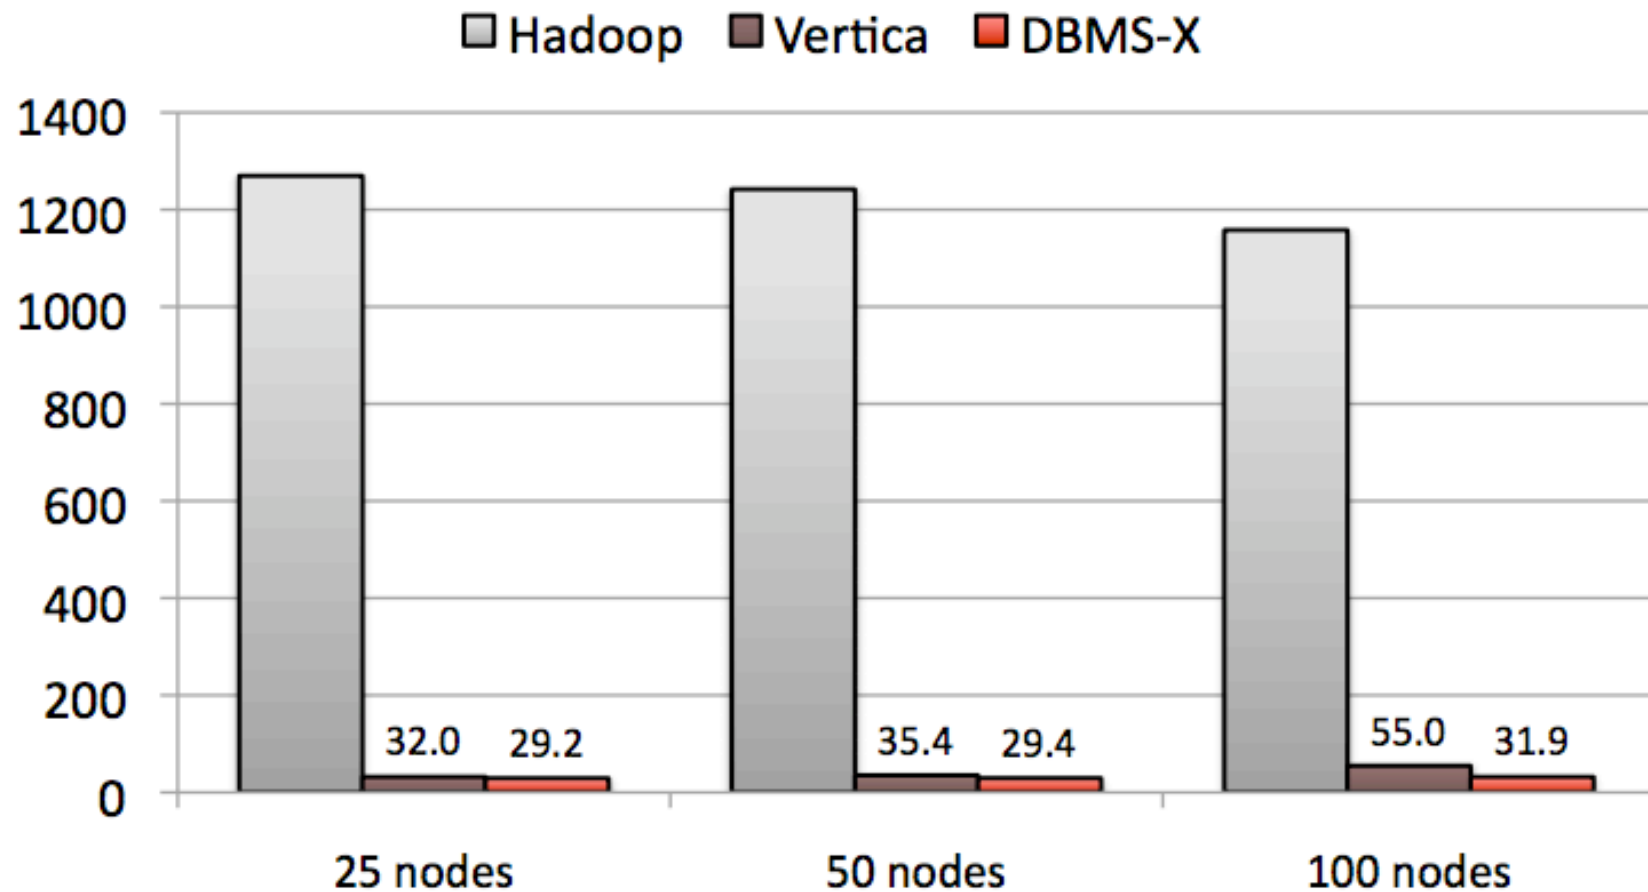
Join Task

- Find the sourceIP that generated the most adRevenue along with its average pageRank.
- Implementations:
 - *DBMSs – Complex SQL using temporary table.*
 - *MapReduce – Three separate MR programs.*

Join Task

```
SELECT INTO TempsourceIP,  
           AVG (pageRank) as avgPageRank,  
           SUM(adRevenue) as totalRevenue  
FROM RankingsAS R  
   , UserVisitsAS UV  
WHERE R.pageURL = UV.destURL  
AND UV.visitDate  
     BETWEEN '2000-01-15'  
     AND '2000-01-22'  
GROUP BY UV.sourceIP;  
  
SELECT sourceIP,  
       totalRevenue,  
       avgPageRank  
FROM Temp  
ORDER BY totalRevenueDESC  
LIMIT 1;
```

Join Task Results



Problems with this analysis?

- Other ways to avoid sequential scans?
- Fault-tolerance in large clusters?
- Tasks that cannot be expressed as queries?

Google's Response: Cluster Size

- Largest known database installations:
 - *Greenplum – 96 nodes – 4.5 PB (eBay) [1]*
 - *Teradata – 72 nodes – 2+ PB (eBay) [1]*
- Largest known MR installations:
 - *Hadoop – 3658 nodes – 1 PB (Yahoo) [2]*
 - *Hive – 600+ nodes – 2.5 PB (Facebook) [3]*

[1] eBay's two enormous data warehouses – April 30th, 2009

<http://www.dbms2.com/2009/04/30/ebays-two-enormous-data-warehouses/>

[2] Hadoop Sorts a Petabyte in 16.25 Hours and a Terabyte in 62 Seconds – May 11th, 2009

http://developer.yahoo.net/blogs/hadoop/2009/05/hadoop_sorts_a_petabyte_in_162.html

[3] Hive - A Petabyte Scale Data Warehouse using Hadoop – June 10th, 2009

http://www.facebook.com/note.php?note_id=89508453919

Concluding Remarks

- What can *MapReduce* learn from *Databases*?
 - *Declarative languages are a good thing.*
 - *Schemas are important.*
- What can *Databases* learn from *MapReduce*?
 - *Query fault-tolerance.*
 - *Support for in situ data.*
 - *Embrace open-source.*

Other Benchmarked Systems

- **HadoopDB (Abadi '09 - Yale)**
 - *Replaced Hadoop filesystem with Postgres.*
 - *Makes JDBC calls inside of MR functions.*
- **Hive (Thusoo '09 - Facebook)**
 - *Data warehouse interface on top of Hadoop.*
 - *Converts SQL-like language to MR programs.*