

1) We need to ensure high availability

2) We also want to support updates

# Example

User: Sue  
Friends: Joe, Kai, ...  
Status: "Headed to new Bond flick"  
Wall: "...", "..."

User: Joe  
Friends: Sue, ...  
Status: "I'm sleepy"  
Wall: "...", "..."

User: Kai  
Friends: Sue, ...  
Status: "Done for tonight"  
Wall: "...", "..."

Write: Update Sue's status. Who sees the new status, and who sees the old one?

Databases: *"Everyone MUST see the same thing, either old or new, no matter how long it takes."*

NoSQL: *"For large applications, we can't afford to wait that long, and maybe it doesn't matter anyway"*

# Example

## Friends

Jim, Sue  
Sue, Jim  
Lin, Joe  
Joe, Lin  
Jim, Kai  
Kai, Jim  
Jim, Lin  
Lin, Jim

## Users

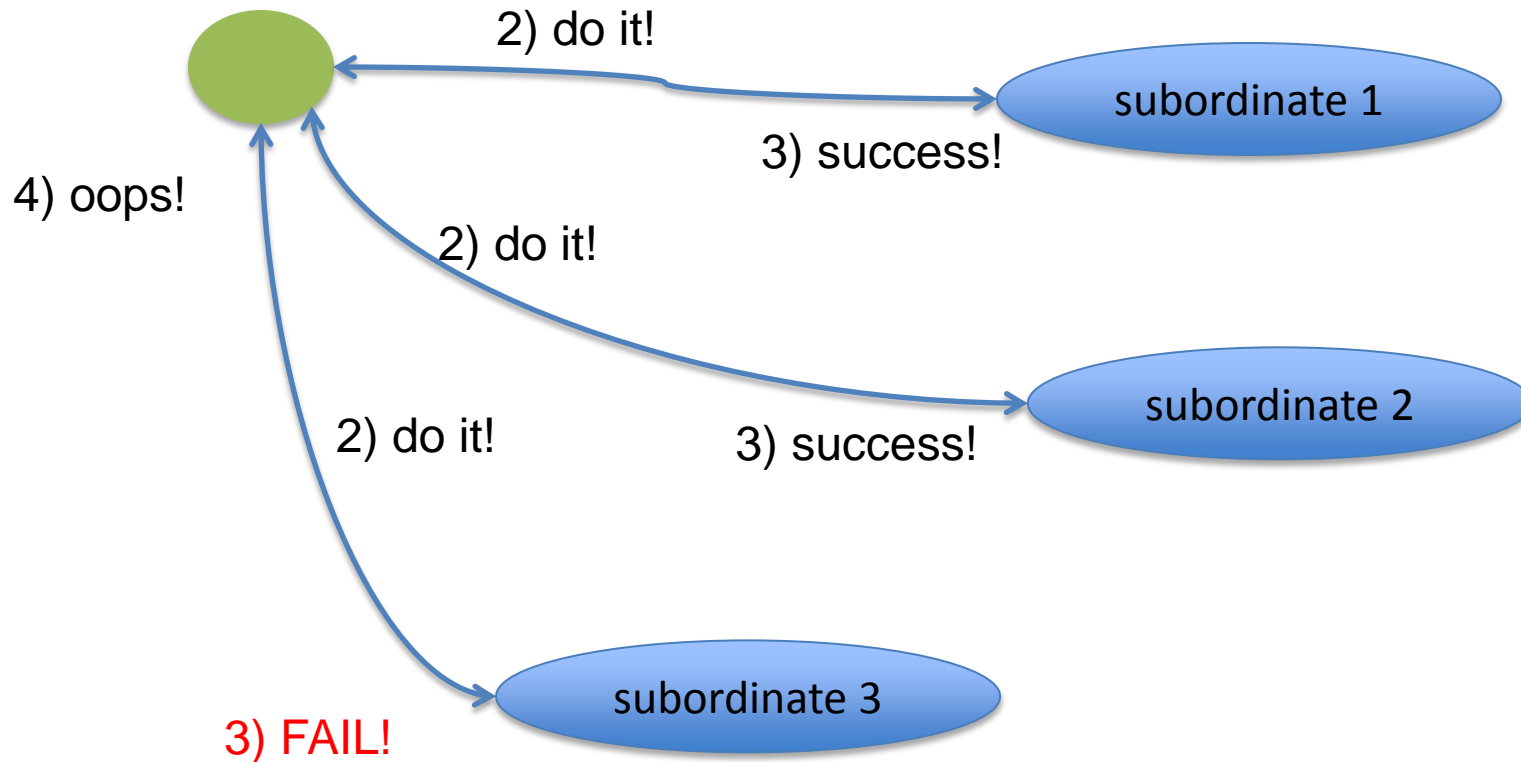
Jim  
Sue  
...

## Posts

Sue: "headed to see new Bond flick"  
Sue: "it was ok"  
Kai: "I'm hungry"

# Two-Phase Commit Motivation

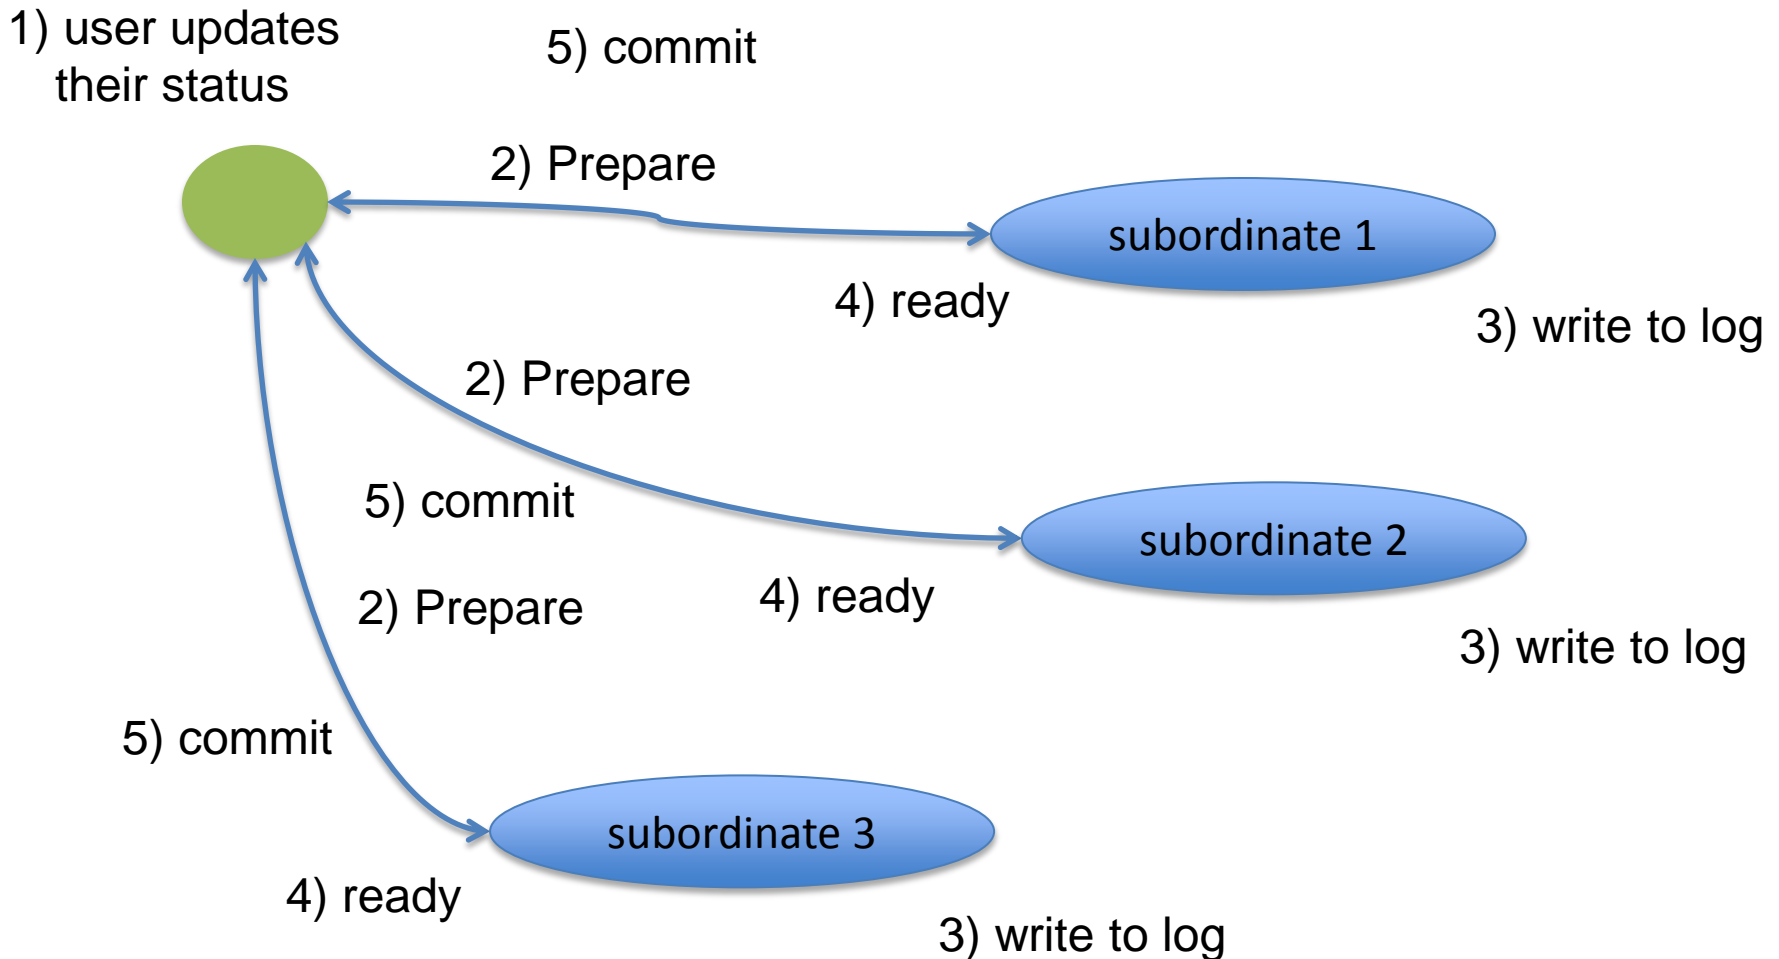
1) user updates  
their status



# Two-Phase Commit

- Phase 1:
  - Coordinator Sends “Prepare to Commit”
  - Subordinates make sure they can do so no matter what
    - Write the action to a log to tolerate failure
  - Subordinates Reply “Ready to Commit”
- Phase 2:
  - If all subordinates ready, send “Commit”
  - If anyone failed, send “Abort”

# Two-Phase Commit



# “Eventual Consistency”

- Write conflicts will eventually propagate throughout the system
  - D. Terry et al., “Managing Update Conflicts in Bayou, a Weakly Connected Replicated Storage System”, SOSP 1995

“We believe that applications must be aware that they may read weakly consistent data and also that their write operations may conflict with those of other users and applications.”

“Moreover, applications must be involved in the detection and resolution of conflicts since these naturally depend on the semantics of the application.”

# Eventual Consistency

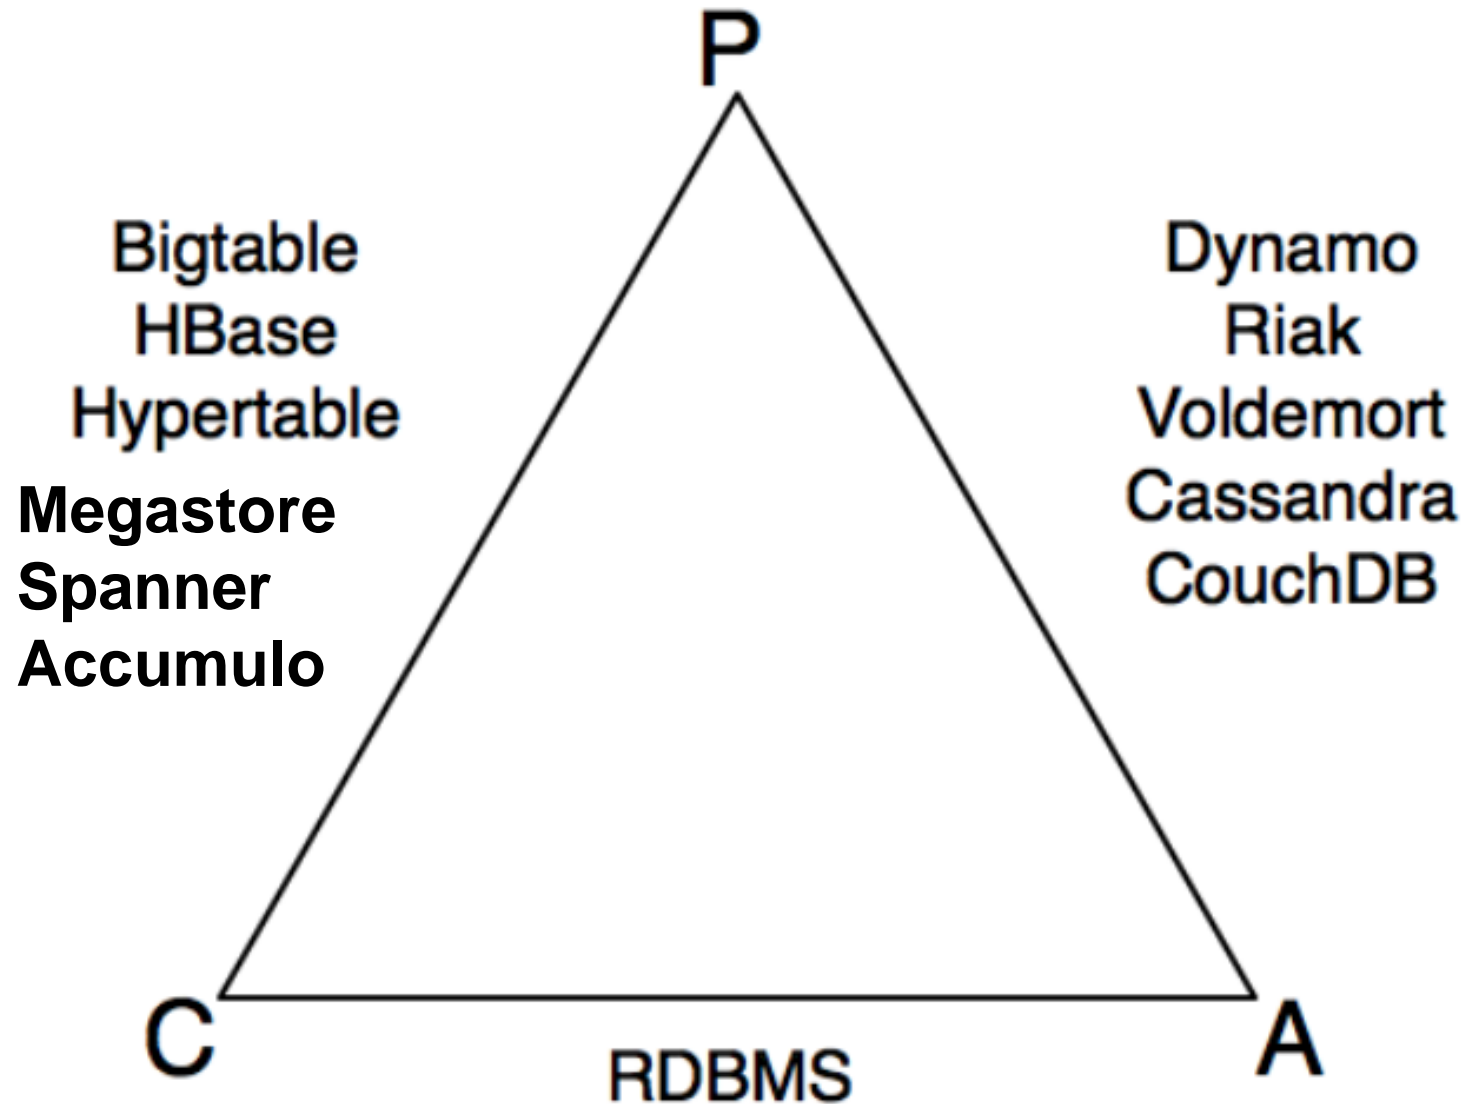
- What the application sees in the meantime is sensitive to replication mechanics and difficult to predict
- Contrast with RDBMS, Paxos: Immediate (or “strong”) consistency, but there may be deadlocks



<i>Year</i>	<i>System/ Paper</i>	<i>Scale to 1000s</i>	<i>Primary Index</i>	<i>Secondary Indexes</i>	<i>Transactions</i>	<i>Joins/ Analytics</i>	<i>Integrity Constraints</i>	<i>Views</i>	<i>Language/ Algebra</i>	<i>Data model</i>	<i>my label</i>
2003	memcached	✓	✓	○	○	○	○	○	○	key-val	nosql
2005	CouchDB	✓	✓	✓	record	MR	○	✓	○	document	nosql
2006	BigTable (Hbase)	✓	✓	✓	record	compat. w/MR	/	○	○	ext. record	nosql
2007	MongoDB	✓	✓	✓	EC, record	○	○	○	○	document	nosql
2007	Dynamo	✓	✓	○	○	○	○	○	○	key-val	nosql
2008	Cassandra	✓	✓	✓	EC, record	○	✓	✓	○	key-val	nosql
2009	Voldemort	✓	✓	○	EC, record	○	○	○	○	key-val	nosql
2009	Riak	✓	✓	✓	EC, record	MR	○			key-val	nosql
2011	Megastore	✓	✓	✓	entity groups	○	/	○	/	tables	nosql
2012	Accumulo	✓	✓	✓	record	compat. w/MR	/	○	○	ext. record	nosql
2012	Spanner	✓	✓	✓	✓	?	✓	✓	✓	tables	sql-like

# CAP Theorem [Brewer 2000, Lynch 2002]

- Consistency
  - Do all applications see all the same data?
- Availability
  - If some nodes fail, does everything still work?
- Partitioning
  - If two sections of your system cannot talk to each other, can they make forward progress on their own?
    - If not, you sacrifice Availability
    - If so, you might have to sacrifice Consistency – can't have everything
- Conventional databases assume no partitioning – clusters were assumed to be small and local
- NoSQL systems may sacrifice consistency



*src: Shashank Tiwari*

Year	System/ Paper	Scale to 1000s	Primary Index	Secondary Indexes	Transactions	Joins/ Analytics	Integrity Constraints	Views	Language/ Algebra	Data model	my label
1971	RDBMS	0	✓	✓	✓	✓	✓	✓	✓	tables	sql-like
2003	memcached	✓	✓	0	0	0	0	0	0	key-val	nosql
2004	MapReduce	✓	0	0	0	✓	0	0	0	key-val	batch
2005	CouchDB	✓	✓	✓	record	MR	0	✓	0	document	nosql
2006	BigTable (Hbase)	✓	✓	✓	record	compat. w/MR	/	0	0	ext. record	nosql
2007	MongoDB	✓	✓	✓	EC, record	0	0	0	0	document	nosql
2007	Dynamo	✓	✓	0	0	0	0	0	0	key-val	nosql
2008	Pig	✓	0	0	0	✓	/	0	✓	tables	sql-like
2008	HIVE	✓	0	0	0	✓	✓	0	✓	tables	sql-like
2008	Cassandra	✓	✓	✓	EC, record	0	✓	✓	0	key-val	nosql
2009	Voldemort	✓	✓	0	EC, record	0	0	0	0	key-val	nosql
2009	Riak	✓	✓	✓	EC, record	MR	0			key-val	nosql
2010	Dremel	✓	0	0	0	/	✓	0	✓	tables	sql-like
2011	Megastore	✓	✓	✓	entity groups	0	/	0	/	tables	nosql
2011	Tenzing	✓	0	0	0	0	✓	✓	✓	tables	sql-like
2011	Spark/Shark	✓	0	0	0	✓	✓	0	✓	tables	sql-like
2012	Spanner	✓	✓	✓	✓	?	✓	✓	✓	tables	sql-like
2012	Accumulo	✓	✓	✓	record	compat. w/MR	/	0	0	ext. record	nosql
2013	Impala	✓	0	0	0	✓	✓	0	✓	tables	sql-like

Rick Cattell's clustering from  
*"Scalable SQL and NoSQL Data Stores"*  
*SIGMOD Record*, 2010

extensible record stores

document stores

key-value stores