

Random Forest Algorithm

Repeat k times:

- Draw a **bootstrap sample** from the dataset
- Train a decision tree
 - Until the tree is maximum size
 - Choose next leaf node
 - Select m attributes at random from the p available
 - Pick the best attribute/split as usual
- Measure **out-of-bag error**
 - Evaluate against the samples that were not selected in the bootstrap
 - Provides measures of strength (inverse error rate), correlation between trees (which increases the forest error rate), and variable importance

Make a prediction by majority vote among the k trees

Random Forests: Variable Importance

- Key Idea: If you scramble the values of a variable and the accuracy of your tree doesn't change much, then the variable isn't very important
- Measure the error increase
- Random Forests are more difficult to interpret than single trees; understanding variable importance helps
 - Ex: Medical applications can't typically rely on black box solutions

Gini Coefficient

- Entropy captured an intuition for “impurity”
 - We want to choose attributes that split records into pure classes
- The gini coefficient measures inequality

$$\text{Gini}(T) = 1 - \sum_{i=1}^n p_i^2$$

Random Forests on Big Data

- Easy to parallelize
 - Trees are built independently
- Handles “small n big p ” problems naturally
 - A subset of attributes are selected by importance

Summary: Decision Trees and Forests

- Representation
 - Decision Trees
 - Sets of decision trees with majority vote
- Evaluation
 - Accuracy
 - Random forests: out-of-bag error
- Optimization
 - Information Gain or Gini Index to measure impurity and select best attributes